

Gary Hyde

Senior Software Engineer, AI Operations Application Winter Garden / Orlando, FL |
gary.robert.hyde@gmail.com | (407) 473-7206 | jobs.garyhyde.com/rapidsos/senior-software-engineer-ai-operations

Dear RapidSOS hiring team,

I am applying for the Senior Software Engineer, AI Operations role, and I want to be direct about why it stood out: the way you describe the work is the way I already work. Full-stack TypeScript and Node.js, building the AI agents, tools, and automations that internal teams actually use, and taking an idea to a deployed production system quickly and independently. I have spent the last several years building products exactly that way, usually as the whole team, because shipping solo leaves no room for handoffs.

The clearest proof is Pathara, an AI product I built end to end. I took it from a blank intake flow through a deployed Claude agent pipeline that does the real research and writing, to a per-candidate results portal with downloads and an application tracker, on Supabase with row-level security and storage throughout. There was no separate backend, frontend, or design resource. The idea, the agent pipeline, the data isolation, and the production deploy were all mine. Lumi and Kora followed the same path, idea to deployed product solo, so going from a rough concept to something running in production without a big team behind me is the normal case for me, not the exception.

Internal AI enablement is the part of this role I am most drawn to, because it is already a habit. On my current team I built code-review tooling that runs several AI models and synthesizes their output into structured scorecards with Ship, Fix, and No-Ship verdicts, which raised our human-in-the-loop review bar. That is the same instinct your AI Operations mandate describes, building the agents and automations that make a team faster, then making sure people trust them enough to use them.

I also live in the tools you name. Claude Code and Cursor are my daily drivers, not things I tried once, and I treat the harness around the model, the tools, the verification steps, the context discipline, as the real source of quality. So when the role asks for someone who can evaluate and adopt agentic coding tools like Claude Code, Cursor, and Copilot and turn them into leverage, that is a job I have already been doing. I am the person who tries a new tool, decides where it earns a place, and wires it into a repeatable workflow the rest of the team can lean on.

A small team with high autonomy is a draw, not a tolerance. I have worked independently and remotely for over ten years, and my default is to scope it, build it, ship it, and write the decision down so others can act on it without a meeting. I would welcome the chance to show how I work in your codebase. Thank you for reading.

Sincerely, Gary Hyde